Package 'RemixAutoML'

March 27, 2019

Title Remix Automated Machine Learning

Version 1.0

Description Automate and ensure high quality output for most of your machine learning and data science tasks. We have high quality functions that run at efficient speed with minimal memory constraints. The library contains functions for supervised learning, unsupervised learning, feature engineering, model evaluation and interpretation, along with some helper functions for graphing.

Depends R (¿= 3.5.0)

SystemRequirements Java (¿= 7.0)

License GPL-2

Encoding UTF-8

Language en-US

URL https://github.com/AdrianAntico/RemixAutoML

BugReports https://github.com/AdrianAntico/RemixAutoML/issues

Contact Adrian Antico

LazyData true

RoxygenNote 6.1.1

Imports data.table, zoo, h2o, lubridate, ggplot2, caTools, forecast, prophet, tsoutliers, stringr, itertools, doParallel, parallel, scatterplot3d, RColorBrewer, grid, monreg, tm, wordcloud, foreach, pROC, doSNOW, recommenderlab

Suggests testthat, sde, knitr, rmarkdown, methods

VignetteBuilder knitr

Author Adrian Antico

Maintainer Adrian Antico <adrianantico@gmail.com>

Date '2019-03-20'

Views MachineLearning,
AutomatedSupervisedLearning,
SupervisedLearning,
AutomatedUnsupervisedLearning,
UnsupervisedLearning,
Clustering,
AnomalyDetection,

FeatureEngineering,
VariableCreation,
ModelEvaluation,
FeatureInterpretation,
VariableInterpretation,
VariableImportance,
AutomatedTimeSeriesForecasting,
TimeSeries,
Forecasting

R **topics documented:**

| AutoH20Modeler | *An Automated Machine Learning Framework using H20* |

Description

Steps in the function include: 1. Logic: Error checking in the modeling arguments from your Construction file 2. ML: Build grid-tuned models and baseline models for comparison and checks which one performs better on validation data 3. Evaluation: Collects the performance metrics for both 4. Evaluation: Generates calibration plots (and boxplots for regression) for the winning model 5. Evaluation: Generates partial dependence calibration plots (and boxplots for regression) for the winning model 6. Evaluation: Generates variable importance tables and a table of non-important features 7. Production: Creates a storage file containing: model name, model path, grid tune performance, baseline performance, and threshold (if classification) and stores that file in your model_path location

Usage

```
AutoH20Modeler(Construct, max_memory, ratios, BL_Trees, nthreads,
  model_path, MaxRuntimeSeconds = 3600, MaxModels = 30,
  TrainData = NULL, TestData = NULL)
```

Arguments

Construct	Core instruction file for automation (see Details below for more information on this)
max_memory	The ceiling amount of memory H20 will utilize
ratios	The percentage of train samples from source data (remainder goes to validation set)
BL_Trees	The number of trees to build in baseline GBM or RandomForest
nthreads	Set the number of threads to run function
model_path	Directory path for where you want your models saved
MaxRuntimeSeconds	
	Number of seconds of run time for grid tuning
MaxModels	Number of models you'd like to have returned
TrainData	Set to NULL or supply a data.table for training data
TestData	Set to NULL or supply a data.table for validation data

Details

The Construct file must be a data.table and the columns need to be in the correct order (see examples). Character columns must be converted to type "Factor". You must remove date columns or convert them to "Factor". For classification models, your target variable needs to be a (0,1) of type "Factor." See the examples below for help with setting up the Construct file for various modeling target variable types. There are examples for regression, classification, multinomial, and quantile regression.

Let's go over the construct file, column by column. The Targets column is where you specify the column number of your target variable (in quotes, e.g. "c(1)"). The Distribution column is where you specify the distribution type for the modeling task. For classification

use bernoulli, for multilabel use multinomial, for quantile use quantile, and for regression, you can choose from the list available in the H20 docs, such as gaussian, poisson, gamma, etc. It's not set up to handle tweedie distributions currently but I can add support if there is demand. The Loss column tells H20 which metric to use for the loss metrics. For regression, I typically use "mse", quantile regression, "mae", classification "auc", and multinomial "logloss". For deeplearning models, you need to use "quadratic", "absolute", and "crossentropy". The Quantile column tells H20 which quantile to use for quantile regression (in decimal form). The ModelName column is the name you wish to give your model as a prefix. The Algorithm column is the model you wish to use: gbm, randomForest, deeplearning, AutoML, XGBoost, LightGBM. The dataName column is the name of your data. The TargetCol column is the column number of your target variable. The FeatureCols column is the column numbers of your features. The CreateDate column is for tracking your model build dates. The GridTune column is a TRUE / FALSE column for whether you want to run a grid tune model for comparison. The ExportValidData column is a TRUE / FALSE column indicating if you want to export the validation data. The ParDep column is where you put the number of partial dependence calibration plots you wish to generate. The PD_Data column is where you specify if you want to generate the partial dependence plots on "All" data, "Validate" data, or "Train" data. The ThreshType column is for classification models. You can specify "f1", "f2", "f0point5", or "CS" for cost sentitive. The FSC column is the feature selection column. Specify the percentage importance cutoff to create a table of "unimportant" features. The tpProfit column is for when you specify "CS" in the ThreshType column. This is your true positive profit. The tnProfit column is for when you specify "CS" in the ThreshType column. This is your true negative profit. The fpProfit column is for when you specify "CS" in the ThreshType column. This is your false positive profit. The fnProfit column is for when you specify "CS" in the ThreshType column. This is your false negative profit. The SaveModel column is a TRUE / FALSE indicator. If you are just testing out models, set this to FALSE. The SaveModelType column is where you specify if you want a "standard" model object saveed or a "mojo" model object saved. The PredsAllData column is a TRUE / FALSE column. Set to TRUE if you want all the predicted values returns (for all data). The TargetEncoding column let's you specify the column number of features you wish to run target encoding on. Set to NA to not run this feature. The SupplyData column lets you supply the data names for training and validation data. Set to NULL if you want the data partitioning to be done internally.

Value

Returns saved models, corrected Construct file, variable importance tables, evaluation and partial dependence calibration plots, model performance measure, etc.

Author(s)

Adrian Antico

See Also

Other Supervised Learning: AutoH20Scoring, AutoNLS, AutoRecommender, AutoTS

Examples

```
## Not run:
# Classification Example
Correl <- 0.85
aa <- data.table::data.table(target = runif(1000))
```

```
aa[, x1 := qnorm(target)]
aa[, x2 := runif(1000)]
aa[, Independent_Variable1 := log(pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable2 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable3 := exp(pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable4 := exp(exp(pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2))))]
aa[, Independent_Variable5 := sqrt(pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable6 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))^0.10]
aa[, Independent_Variable7 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))^0.25]
aa[, Independent_Variable8 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))^0.75]
aa[, Independent_Variable9 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))^2]
aa[, Independent_Variable10 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))^4]
aa[, ':=' (x1 = NULL, x2 = NULL)]
aa[, target := as.factor(ifelse(target > 0.5,1,0))]
Construct <- data.table::data.table(Targets = rep("target",3),
                            Distribution    = c("bernoulli",
                                                "bernoulli",
                                                "bernoulli"),
                            Loss            = c("AUC","AUC","CrossEntropy"),
                            Quantile        = rep(NA,3),
                            ModelName       = c("GBM","DRF","DL"),
                            Algorithm       = c("gbm",
                                                "randomForest",
                                                "deeplearning"),
                            dataName        = rep("aa",3),
                            TargetCol       = rep(c("1"),3),
                            FeatureCols     = rep(c("2:11"),3),
                            CreateDate      = rep(Sys.time(),3),
                            GridTune        = rep(FALSE,3),
                            ExportValidData = rep(TRUE,3),
                            ParDep          = rep(2,3),
                            PD_Data         = rep("All",3),
                            ThreshType      = rep("f1",3),
                            FSC             = rep(0.001,3),
                            tpProfit        = rep(NA,3),
                            tnProfit        = rep(NA,3),
                            fpProfit        = rep(NA,3),
                            fnProfit        = rep(NA,3),
                            SaveModel       = rep(FALSE,3),
                            SaveModelType   = c("Mojo","standard","mojo"),
                            PredsAllData    = rep(TRUE,3),
                            TargetEncoding  = rep(NA,3),
                            SupplyData      = rep(FALSE,3))
AutoH20Modeler(Construct,
                max_memory = "28G",
                ratios = 0.75,
                BL_Trees = 500,
```

```
                nthreads = 5,
                model_path = getwd(),
                MaxRuntimeSeconds = 3600,
                MaxModels = 30,
                TrainData = NULL,
                TestData  = NULL)

# Multinomial Example
Correl <- 0.85
aa <- data.table::data.table(target = runif(1000))
aa[, x1 := qnorm(target)]
aa[, x2 := runif(1000)]
aa[, Independent_Variable1 := log(pnorm(Correl * x1 +
                                  sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable2 := (pnorm(Correl * x1 +
                                  sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable3 := exp(pnorm(Correl * x1 +
                                  sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable4 := exp(exp(pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2))))]
aa[, Independent_Variable5 := sqrt(pnorm(Correl * x1 +
                                  sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable6 := (pnorm(Correl * x1 +
                                sqrt(1-Correl^2) * qnorm(x2)))^0.10]
aa[, Independent_Variable7 := (pnorm(Correl * x1 +
                                sqrt(1-Correl^2) * qnorm(x2)))^0.25]
aa[, Independent_Variable8 := (pnorm(Correl * x1 +
                                sqrt(1-Correl^2) * qnorm(x2)))^0.75]
aa[, Independent_Variable9 := (pnorm(Correl * x1 +
                                sqrt(1-Correl^2) * qnorm(x2)))^2]
aa[, Independent_Variable10 := (pnorm(Correl * x1 +
                                sqrt(1-Correl^2) * qnorm(x2)))^4]
aa[, ':=' (x1 = NULL, x2 = NULL)]
aa[, target := as.factor(ifelse(target < 0.33,"A",ifelse(target < 0.66, "B","C")))]
Construct <- data.table::data.table(Targets = rep("target",3),
                            Distribution    = c("multinomial",
                                                "multinomial",
                                                "multinomial"),
                            Loss            = c("auc","logloss","accuracy"),
                            Quantile        = rep(NA,3),
                            ModelName       = c("GBM","DRF","DL"),
                            Algorithm       = c("gbm",
                                                "randomForest",
                                                "deeplearning"),
                            dataName        = rep("aa",3),
                            TargetCol       = rep(c("1"),3),
                            FeatureCols     = rep(c("2:11"),3),
                            CreateDate      = rep(Sys.time(),3),
                            GridTune        = rep(FALSE,3),
                            ExportValidData = rep(TRUE,3),
                            ParDep          = rep(NA,3),
                            PD_Data         = rep("All",3),
                            ThreshType      = rep("f1",3),
                            FSC             = rep(0.001,3),
                            tpProfit        = rep(NA,3),
                            tnProfit        = rep(NA,3),
                            fpProfit        = rep(NA,3),
```

```
                                          fnProfit        = rep(NA,3),
                                          SaveModel       = rep(FALSE,3),
                                          SaveModelType   = c("Mojo","standard","mojo"),
                                          PredsAllData    = rep(TRUE,3),
                                          TargetEncoding  = rep(NA,3),
                                          SupplyData      = rep(FALSE,3))

AutoH20Modeler(Construct,
               max_memory = "28G",
               ratios = 0.75,
               BL_Trees = 500,
               nthreads = 5,
               model_path = getwd(),
               MaxRuntimeSeconds = 3600,
               MaxModels = 30,
               TrainData = NULL,
               TestData  = NULL)

# Regression Example
Correl <- 0.85
aa <- data.table::data.table(target = runif(1000))
aa[, x1 := qnorm(target)]
aa[, x2 := runif(1000)]
aa[, Independent_Variable1 := log(pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable2 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable3 := exp(pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable4 := exp(exp(pnorm(Correl * x1 +
                                      sqrt(1-Correl^2) * qnorm(x2))))]
aa[, Independent_Variable5 := sqrt(pnorm(Correl * x1 +
                                     sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable6 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))^0.10]
aa[, Independent_Variable7 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))^0.25]
aa[, Independent_Variable8 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))^0.75]
aa[, Independent_Variable9 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))^2]
aa[, Independent_Variable10 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))^4]
aa[, ':=' (x1 = NULL, x2 = NULL)]
Construct <- data.table::data.table(Targets = rep("target",3),
                                    Distribution    = c("gaussian",
                                                        "gaussian",
                                                        "gaussian"),
                                    Loss            = c("MSE","MSE","Quadratic"),
                                    Quantile        = rep(NA,3),
                                    ModelName       = c("GBM","DRF","DL"),
                                    Algorithm       = c("gbm",
                                                        "randomForest",
                                                        "deeplearning"),
                                    dataName        = rep("aa",3),
                                    TargetCol       = rep(c("1"),3),
                                    FeatureCols     = rep(c("2:11"),3),
```

```
                                                      CreateDate      = rep(Sys.time(),3),
                                                      GridTune        = rep(FALSE,3),
                                                      ExportValidData = rep(TRUE,3),
                                                      ParDep          = rep(2,3),
                                                      PD_Data         = rep("All",3),
                                                      ThreshType      = rep("f1",3),
                                                      FSC             = rep(0.001,3),
                                                      tpProfit        = rep(NA,3),
                                                      tnProfit        = rep(NA,3),
                                                      fpProfit        = rep(NA,3),
                                                      fnProfit        = rep(NA,3),
                                                      SaveModel       = rep(FALSE,3),
                                                      SaveModelType   = c("Mojo","standard","mojo"),
                                                      PredsAllData    = rep(TRUE,3),
                                                      TargetEncoding  = rep(NA,3),
                                                      SupplyData      = rep(FALSE,3))
AutoH20Modeler(Construct,
               max_memory = "28G",
               ratios = 0.75,
               BL_Trees = 500,
               nthreads = 5,
               model_path = getwd(),
               MaxRuntimeSeconds = 3600,
               MaxModels = 30,
               TrainData = NULL,
               TestData  = NULL)


# Quantile Regression Example
Correl <- 0.85
aa <- data.table::data.table(target = runif(1000))
aa[, x1 := qnorm(target)]
aa[, x2 := runif(1000)]
aa[, Independent_Variable1 := log(pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable2 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable3 := exp(pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable4 := exp(exp(pnorm(Correl * x1 +
                                        sqrt(1-Correl^2) * qnorm(x2))))]
aa[, Independent_Variable5 := sqrt(pnorm(Correl * x1 +
                                     sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable6 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))^0.10]
aa[, Independent_Variable7 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))^0.25]
aa[, Independent_Variable8 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))^0.75]
aa[, Independent_Variable9 := (pnorm(Correl * x1 +
                                 sqrt(1-Correl^2) * qnorm(x2)))^2]
aa[, Independent_Variable10 := (pnorm(Correl * x1 +
                                  sqrt(1-Correl^2) * qnorm(x2)))^4]
aa[, ':=' (x1 = NULL, x2 = NULL)]
Construct <- data.table::data.table(Targets = rep("target",3),
                                    Distribution    = c("quantile",
                                                        "quantile"),
                                    Loss            = c("MAE","Absolute"),
```

```
                                        Quantile        = rep(0.75,2),
                                        ModelName       = c("GBM","DL"),
                                        Algorithm       = c("gbm",
                                                            "deeplearning"),
                                        dataName        = rep("aa",2),
                                        TargetCol       = rep(c("1"),2),
                                        FeatureCols     = rep(c("2:11"),2),
                                        CreateDate      = rep(Sys.time(),2),
                                        GridTune        = rep(FALSE,2),
                                        ExportValidData = rep(TRUE,2),
                                        ParDep          = rep(4,2),
                                        PD_Data         = rep("All",2),
                                        ThreshType      = rep("f1",2),
                                        FSC             = rep(0.001,2),
                                        tpProfit        = rep(NA,2),
                                        tnProfit        = rep(NA,2),
                                        fpProfit        = rep(NA,2),
                                        fnProfit        = rep(NA,2),
                                        SaveModel       = rep(FALSE,2),
                                        SaveModelType   = c("Mojo","mojo"),
                                        PredsAllData    = rep(TRUE,2),
                                        TargetEncoding  = rep(NA,2),
                                        SupplyData      = rep(FALSE,2))
AutoH20Modeler(Construct,
               max_memory = "28G",
               ratios = 0.75,
               BL_Trees = 500,
               nthreads = 5,
               model_path = getwd(),
               MaxRuntimeSeconds = 3600,
               MaxModels = 30,
               TrainData = NULL,
               TestData  = NULL)

## End(Not run)
```

AutoH20Scoring   *AutoH20Scoring is the complement of AutoH20Modeler.*

Description

AutoH20Scoring is the complement of AutoH20Modeler. Use this for scoring models. You can score regression, quantile regression, classification, multinomial, and text models (built with the Word2VecModel function). You can also use this to score multioutcome models so long as the there are two models: one for predicting the count of outcomes (a count outcome in character form) and a multinomial model on the label data. You will want to ensure you have a record for each label in your training data in (0,1) as factor form.

Usage

```
AutoH20Scoring(Features = data, GridTuneRow = c(1:3),
  ScoreMethod = "Standard", TargetType = rep("multinomial", 3),
  ClassVals = rep("probs", 3), NThreads = 6, MaxMem = "28G",
  JavaOptions = "-Xmx1g -XX:ReservedCodeCacheSize=256m",
  FilesPath = getwd(), H20ShutDown = rep(FALSE, 3))
```

Arguments

Features	This is a data.table of features for scoring.
GridTuneRow	Numeric. The row numbers of grid_tuned_paths or StoreFile containing the model you wish to score
ScoreMethod	"Standard" or "Mojo"
TargetType	"Regression", "Classification", "Multinomial", "Text", "MultiOutcome". MultiOutcome must be two multinomial models, a count model (the count of outcomes, as a character value), and the multinomial model predicting the labels.
ClassVals	Choose from "p1", "Probs", "Label", or "All"
NThreads	Number of available threads for H20
MaxMem	Amount of memory to dedicate to H20
JavaOptions	Modify to your machine if the default doesn't work
FilesPath	Set this to the folder where your models are saved (and hence where your grid_tuned_paths.Rdata file resides)
H20ShutDown	TRUE to shutdown H20 after the run (do this if you are scoring once and not after that for a long time). Use FALSE if you will be repeatedly scoring and shutdown somewhere else in your script

Value

Returns a list of predicted values. Each list element contains the predicted values from a single model predict call.

Author(s)

Adrian Antico

See Also

Other Supervised Learning: AutoH20Modeler, AutoNLS, AutoRecommender, AutoTS

Examples

```
## Not run:
# Multinomial Example
Correl <- 0.85
aa <- data.table::data.table(target = runif(1000))
aa[, x1 := qnorm(target)]
aa[, x2 := runif(1000)]
aa[, Independent_Variable1 := log(pnorm(Correl * x1 +
                                          sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable2 := (pnorm(Correl * x1 +
                                       sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable3 := exp(pnorm(Correl * x1 +
                                          sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable4 := exp(exp(pnorm(Correl * x1 +
                                              sqrt(1-Correl^2) * qnorm(x2))))]
aa[, Independent_Variable5 := sqrt(pnorm(Correl * x1 +
                                           sqrt(1-Correl^2) * qnorm(x2)))]
aa[, Independent_Variable6 := (pnorm(Correl * x1 +
                                       sqrt(1-Correl^2) * qnorm(x2)))^0.10]
```

```
aa[, Independent_Variable7 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))^0.25]
aa[, Independent_Variable8 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))^0.75]
aa[, Independent_Variable9 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))^2]
aa[, Independent_Variable10 := (pnorm(Correl * x1 +
                                    sqrt(1-Correl^2) * qnorm(x2)))^4]
aa[, ':=' (x1 = NULL, x2 = NULL)]
aa[, target := as.factor(ifelse(target < 0.33,"A",ifelse(target < 0.66, "B","C")))]
Construct <- data.table::data.table(Targets = rep("target",3),
                              Distribution    = c("multinomial",
                                                  "multinomial",
                                                  "multinomial"),
                        Loss            = c("logloss","logloss","CrossEntropy"),
                        Quantile        = rep(NA,3),
                        ModelName       = c("GBM","DRF","DL"),
                        Algorithm       = c("gbm",
                                            "randomForest",
                                            "deeplearning"),
                        dataName        = rep("aa",3),
                        TargetCol       = rep(c("1"),3),
                        FeatureCols     = rep(c("2:11"),3),
                        CreateDate      = rep(Sys.time(),3),
                        GridTune        = rep(FALSE,3),
                        ExportValidData = rep(TRUE,3),
                        ParDep          = rep(NA,3),
                        PD_Data         = rep("All",3),
                        ThreshType      = rep("f1",3),
                        FSC             = rep(0.001,3),
                        tpProfit        = rep(NA,3),
                        tnProfit        = rep(NA,3),
                        fpProfit        = rep(NA,3),
                        fnProfit        = rep(NA,3),
                        SaveModel       = rep(FALSE,3),
                        SaveModelType   = c("Mojo","mojo","mojo"),
                        PredsAllData    = rep(TRUE,3),
                        TargetEncoding  = rep(NA,3),
                        SupplyData      = rep(FALSE,3))

AutoH20Modeler(Construct,
               max_memory = "28G",
               ratios = 0.75,
               BL_Trees = 500,
               nthreads = 5,
               model_path = getwd(),
               MaxRuntimeSeconds = 3600,
               MaxModels = 30,
               TrainData = NULL,
               TestData  = NULL)

N <- 3
data <- AutoH20Scoring(Features     = data,
                       GridTuneRow  = c(1:N),
                       ScoreMethod  = "standard",
                       TargetType   = rep("multinomial",N),
                       ClassVals    = rep("Probs",N),
```

```
                        NThreads    = 6,
                        MaxMem      = "28G",
                        JavaOptions = '-Xmx1g -XX:ReservedCodeCacheSize=256m',
                        FilesPath   = getwd(),
                        H20ShutDown = rep(FALSE,N))

## End(Not run)
```

AutoH20TextPrepScoring

AutoH20TextPrepScoring is for NLP scoring

Description

This function returns prepared tokenized data for H20 Word2VecModeler scoring

Usage

```
AutoH20TextPrepScoring(data, string, MaxMem, NThreads)
```

Arguments

data	The text data
string	The name of the string column to prepare
MaxMem	Amount of memory you want to let H20 utilize
NThreads	The number of threads you want to let H20 utilize

Author(s)

Adrian Antico

See Also

Other Misc: AutoWordFreq, ChartTheme, PrintObjectsSize, RecomDataCreate, RemixTheme, SimpleCap, multiplot, percRank, tempDatesFun, tokenizeH20

AutoKMeans   *AutoKMeans Automated row clustering for mixed column types*

Description

AutoKMeans adds a column to your original data with a cluster number identifier. Uses glrm (grid tune-able) and then k-means to find optimal k.

Usage

```
AutoKMeans(data, GridTuneGLRM = TRUE, GridTuneKMeans = TRUE,
  nthreads = 4, MaxMem = "14G", glrmCols = 3:ncol(data),
  IgnoreConstCols = TRUE, glrmFactors = 5, Loss = "Absolute",
  glrmMaxIters = 1000, SVDMethod = "Randomized",
  MaxRunTimeSecs = 3600, KMeansK = 50, KMeansMetric = "totss")
```

Arguments

`data`	is the source time series data.table
`GridTuneGLRM`	If you want to grid tune the glrm model, set to TRUE, FALSE otherwise
`GridTuneKMeans`	If you want to grid tuen the KMeans model, set to TRUE, FALSE otherwise
`nthreads`	set based on number of threads your machine has available
`MaxMem`	set based on the amount of memory your machine has available
`glrmCols`	the column numbers for the glrm
`IgnoreConstCols`	
	tell H20 to ignore any columns that have zero variance
`glrmFactors`	similar to the number of factors to return from PCA
`Loss`	set to one of "Quadratic", "Absolute", "Huber", "Poisson", "Hinge", "Logistic", "Periodic"
`glrmMaxIters`	max number of iterations
`SVDMethod`	choose from "Randomized","GramSVD","Power"
`MaxRunTimeSecs`	set the timeout for max run time
`KMeansK`	number of factors to test out in k-means to find the optimal number
`KMeansMetric`	pick the metric to identify top model in grid tune c("totss","betweenss","withinss")

Value

Original data.table with added column with cluster number identifier

Author(s)

Adrian Antico

See Also

Other Unsupervised Learning: GenTSAnomVars, ResidualOutliers

Examples

```
## Not run:
data <- data.table::as.data.table(iris)
data <- AutoKMeans(data,
                GridTuneGLRM = TRUE,
                GridTuneKMeans = TRUE,
                nthreads = 8,
                MaxMem = "28G",
                glrmCols = 1:(ncol(data)-1),
                IgnoreConstCols = TRUE,
                glrmFactors = 2,
                Loss = "Absolute",
                glrmMaxIters = 1000,
                SVDMethod = "Randomized",
                MaxRunTimeSecs = 3600,
                KMeansK = 5)
unique(data[["Species"]])
unique(data[["ClusterID"]])
temp <- data[, mean(ClusterID), by = "Species"]
```

```
Setosa <- round(temp[Species == "setosa", V1][[1]],0)
Versicolor <- round(temp[Species == "versicolor", V1][[1]],0)
Virginica <- round(temp[Species == "virginica", V1][[1]],0)
data[, Check := "a"]
data[ClusterID == eval(Setosa), Check := "setosa"]
data[ClusterID == eval(Virginica), Check := "virginica"]
data[ClusterID == eval(Versicolor), Check := "versicolor"]
data[, Acc := as.numeric(ifelse(Check == Species, 1, 0))]
data[, mean(Acc)][[1]]

## End(Not run)
```

AutoNLS   *AutoNLS is a function for automatically building nls models*

Description

This function will build models for 9 different nls models, along with a non-parametric monotonic regression and a polynomial regression. The models are evaluated, a winner is picked, and the predicted values are stored in your data table.

Usage

```
AutoNLS(data, y, x, monotonic = TRUE)
```

Arguments

data	Data is the data table you are building the modeling on
y	Y is the target variable name in quotes
x	X is the independent variable name in quotes
monotonic	This is a TRUE/FALSE indicator - choose TRUE if you want monotonic regression over polynomial regression

Value

A list containing 1: A data table with your original column replaced by the nls model predictions; 2: The model name; 3: The winning model to later use; 4: Model metrics for models with ability to build.

Author(s)

Adrian Antico

See Also

Other Supervised Learning: AutoH20Modeler, AutoH20Scoring, AutoRecommender, AutoTS

Examples

```
# Create Growth Data
data <-
  data.table::data.table(Target = seq(1, 500, 1),
                         Variable = rep(1, 500))
for (i in as.integer(1:500)) {
  if (i == 1) {
    var <- data[i, "Target"][[1]]
    data.table::set(data,
                    i = i,
                    j = 2L,
                    value = var * (1 + runif(1) / 100))
  } else {
    var <- data[i - 1, "Variable"][[1]]
    data.table::set(data,
                    i = i,
                    j = 2L,
                    value = var * (1 + runif(1) / 100))
  }
}

# Add jitter to Target
data[, Target := jitter(Target,
                        factor = 0.25)]

# To keep original values
data1 <- data.table::copy(data)

# Merge and Model data
data11 <- AutoNLS(
  data = data,
  y = "Target",
  x = "Variable",
  monotonic = TRUE
)

# Join predictions to source data
data2 <- merge(
  data1,
  data11[[1]],
  by = "Variable",
  all = FALSE
)

# Plot output
ggplot2::ggplot(data2, ggplot2::aes(x = Variable)) +
  ggplot2::geom_line(ggplot2::aes(y = data2[["Target.x"]],
                                  color = "Target")) +
  ggplot2::geom_line(ggplot2::aes(y = data2[["Target.y"]],
                                  color = "Predicted")) +
 RemixAutoML::ChartTheme(Size = 12) +
  ggplot2::ggtitle(paste0("Growth Models AutoNLS: ",
                          data11[[2]])) +
  ggplot2::ylab("Target Variable") +
  ggplot2::xlab("Independent Variable") +
  ggplot2::scale_colour_manual("Values",
```

```
                                    breaks = c("Target",
                                               "Predicted"),
                              values = c("red",
                                         "blue"))
summary(data11[[3]])
data11[[4]]
```

AutoRecommender	*Automatically build the best recommendere model.*

Description

This function returns the winning model

Usage

```
AutoRecommender(data, Partition = "Split", KFolds = 2, Ratio = 0.75,
  RatingType = "TopN", RatingsKeep = 20,
  SkipModels = "AssociationRules")
```

Arguments

data	This is your BinaryRatingsMatrix. See function RecomDataCreate
Partition	Choose from "split", "cross-validation", "bootstrap". See evaluationScheme in recommenderlab for details.
KFolds	Choose 2 for traditional train and test. Choose greater than 2 for the number of cross validations
Ratio	The ratio for train and test. E.g. 0.75 for 75 percent data allocated to training
RatingType	Choose from "topNList", "ratings", "ratingMatrix"
RatingsKeep	The total ratings you wish to return. Default is 20.
SkipModels	AssociationRules runs the slowest and may crash your system. Choose from: "AssociationRules","ItemBasedCF","UserBasedCF","PopularItems","RandomItems"

Value

The winning model

Author(s)

Adrian Antico and Douglas Pestana

See Also

Other Supervised Learning: AutoH20Modeler, AutoH20Scoring, AutoNLS, AutoTS

Examples

```
## Not run:
ModelResults <- AutoRecommender(BinaryRatingsMatrix,
                                Partition = "Split",
                                KFolds = 2,
                                Ratio = 0.75,
                                RatingType = "TopN",
                                RatingsKeep = 20,
                                SkipModels = "AssociationRules")

## End(Not run)
```

AutoTS   *AutoTS is an automated time series modeling function*

Description

AutoTS builds the best time series models for each type, compares all types, selects the winner, and generates a forecast.

Usage

```
AutoTS(data, TargetName = "Targets", DateName = "DateTime",
  FCPeriods = 30, HoldOutPeriods = 30, TimeUnit = "day", Lags = 25,
  SLags = 2, NumCores = 4, SkipModels = NULL, StepWise = TRUE)
```

Arguments

data	is the source time series data.table
TargetName	is the name of the dependent variable in your data.table
DateName	is the name of the date column in your data.table
FCPeriods	is the number of periods into the future you wish to forecast
HoldOutPeriods	is the number of periods to use for validation testing
TimeUnit	is the level of aggregation your dataset comes in
Lags	is the number of lags you wish to test in various models (same with moving averages)
SLags	is the number of seasonal lags you wish to test in various models (same with moving averages)
NumCores	is the number of cores available on your computer
SkipModels	Don't run specified models - e.g. exclude all models "ARFIMA" "ARIMA" "ETS" "NNET" "TBATS" "TSLM" "PROPHET"
StepWise	Set to TRUE to have ARIMA and ARFIMA run a stepwise selection process. Otherwise, all models will be generated in parallel execution, but still run much slower.

Value

Returns a list containing 1: A data.table object with a date column and the forecasted values; 2: The model evaluation results; 3: The winning model for later use if desired.

Author(s)

Adrian Antico and Douglas Pestana

See Also

Other Supervised Learning: AutoH20Modeler, AutoH20Scoring, AutoNLS, AutoRecommender

Examples

```
data <- data.table::data.table(DateTime = as.Date(Sys.time()),
  Target = stats::filter(rnorm(1000,
                              mean = 50,
                              sd = 20),
                      filter=rep(1,10),
                      circular=TRUE))
data[, temp := seq(1:1000)][, DateTime := DateTime - temp][, temp := NULL]
data <- data[order(DateTime)]
output <-    AutoTS(data,
                    TargetName    = "Target",
                    DateName      = "DateTime",
                    FCPeriods     = 30,
                    HoldOutPeriods = 30,
                    TimeUnit      = "day",
                    Lags          = 5,
                    SLags         = 1,
                    NumCores      = 4,
                    SkipModels    = NULL,
                    StepWise      = TRUE)
ForecastData <- output[[1]]
ModelEval    <- output[[2]]
WinningModel <- output[[3]]
```

AutoWord2VecModeler   *Automated word2vec data generation via H20*

Description

This function allows you to automatically build a word2vec model and merge the data onto your supplied dataset

Usage

```
AutoWord2VecModeler(data, stringCol = c("Text_Col1", "Text_Col2"),
  KeepStringCol = FALSE, model_path = getwd(), vects = 100,
  SaveStopWords = FALSE, MinWords = 1, WindowSize = 12,
  Epochs = 25, StopWords = NULL, SaveModel = "standard",
  Threads = 6, MaxMemory = "28G")
```

Arguments

data	Source data table to merge vects onto
stringCol	A string name for the column to convert via word2vec

KeepStringCol	Set to TRUE if you want to keep the original string column that you convert via word2vec
model_path	A string path to the location where you want the model and metadata stored
vects	The number of vectors to retain from the word2vec model
SaveStopWords	Set to TRUE to save the stop words used
MinWords	For H20 word2vec model
WindowSize	For H20 word2vec model
Epochs	For H20 word2vec model
StopWords	For H20 word2vec model
SaveModel	Set to "standard" to save normally; set to "mojo" to save as mojo. NOTE: while you can save a mojo, I haven't figured out how to score it in the AutoH20Scoring function.
Threads	Number of available threads you want to dedicate to model building
MaxMemory	Amount of memory you want to dedicate to model building

Author(s)

Adrian Antico

See Also

Other Feature Engineering: DT_GDL_Feature_Engineering, DummifyDT, FAST_GDL_Feature_Engineering, GDL_Feature_Engineering, ModelDataPrep, Scoring_GDL_Feature_Engineering

Examples

```
## Not run:
data <- Word2VecModel(data,
                      stringCol     = c("Text_Col1",
                                         "Text_Col2"),
                      KeepStringCol = FALSE,
                      model_path    = getwd(),
                      vects         = 5,
                      SaveStopWords = FALSE,
                      MinWords      = 1,
                      WindowSize    = 1,
                      Epochs        = 25,
                      StopWords     = NULL,
                      SaveModel     = "standard",
                      Threads       = 6,
                      MaxMemory     = "28G")

## End(Not run)
```

AutoWordFreq   *Automated Word Frequency and Word Cloud Creation*

Description

This function builds a word frequency table and a word cloud. It prepares data, cleans text, and generates output.

Usage

```
AutoWordFreq(data, TextColName = "DESCR",
  ClusterCol = "ClusterAllNoTarget", ClusterID = 0,
  RemoveEnglishStopwords = TRUE, Stemming = TRUE,
  StopWords = c("bla", "blab2"))
```

Arguments

data	Source data table
TextColName	A string name for the column
ClusterCol	Set to NULL to ignore, otherwise set to Cluster column name (or factor column name)
ClusterID	Must be set if ClusterCol is defined. Set to cluster ID (or factor level)
RemoveEnglishStopwords	
	Set to TRUE to remove English stop words, FALSE to ignore
Stemming	Set to TRUE to run stemming on your text data
StopWords	Add your own stopwords, in vector format

Author(s)

Adrian Antico

See Also

Other Misc: AutoH20TextPrepScoring, ChartTheme, PrintObjectsSize, RecomDataCreate, RemixTheme, SimpleCap, multiplot, percRank, tempDatesFun, tokenizeH20

Examples

```
## Not run:
data <- WordFreq(data,
               TextColName = "DESCR",
               ClusterCol = "ClusterAllNoTarget",
               ClusterID = 0,
               RemoveEnglishStopwords = TRUE,
               Stemming = TRUE,
               StopWords = c("bla1", "bla2")

## End(Not run)
```

ChartTheme   *ChartTheme function is a ggplot theme generator for ggplots*

Description

This function helps your ggplots look professional with the choice of the two main colors that will dominate the theme

Usage

```
ChartTheme(Size = 12)
```

Arguments

Size   The size of the axis labels and title

Value

An object to pass along to ggplot objects following the "+" sign

Author(s)

Adrian Antico

See Also

Other Misc: AutoH20TextPrepScoring, AutoWordFreq, PrintObjectsSize, RecomDataCreate, RemixTheme, SimpleCap, multiplot, percRank, tempDatesFun, tokenizeH20

Examples

```
data <- data.table::data.table(DateTime = as.Date(Sys.time()),
  Target = stats::filter(rnorm(1000,
                               mean = 50,
                               sd = 20),
                         filter=rep(1,10),
                         circular=TRUE))
data[, temp := seq(1:1000)][, DateTime := DateTime - temp][, temp := NULL]
data <- data[order(DateTime)]
p <- ggplot2::ggplot(data, ggplot2::aes(x = DateTime, y = Target)) + ggplot2::geom_line()
p <- p + ChartTheme(Size = 12)
```

```
DT_GDL_Feature_Engineering
```
              An Automated Feature Engineering Function Using data.table
              frollmean

Description

Builds autoregressive and moving average from target columns and distributed lags and distributed moving average for independent features distributed across time. On top of that, you can also create time between instances along with their associated lags and moving averages. This function works for data with groups and without groups.

Usage

```
DT_GDL_Feature_Engineering(data, lags = c(seq(1, 50, 1)),
  periods = c(seq(5, 95, 5)), statsNames = c("MA"),
  targets = c("qty"), groupingVars = c("Group1", "Group2"),
  sortDateName = c("date"), timeDiffTarget = c("TimeDiffName"),
  timeAgg = c("days"), WindowingLag = 0, Type = c("Lag"),
  Timer = TRUE, SkipCols = NULL, SimpleImpute = TRUE)
```

Arguments

data	The data source you want to run the function on
lags	The list of specific lags you want to have generated
periods	The number of periods for the rolling stats
statsNames	The corresponding names to append to your colnames created associated with statsFuns
targets	The column(s) in which you will build your lags and rolling stats
groupingVars	Categorical variables you will build your lags and rolling stats by
sortDateName	String name of your core date column in your transaction data
timeDiffTarget	List a name in order to create time between events with assiciated lags and rolling features
timeAgg	Unit of time to aggregate by
WindowingLag	Build moving stats off of target column(s) or one of their lags (1+)
Type	input "Lag" if you want features built on historical values; use "Lead" if you want features built on future values
Timer	Set to TRUE if you want a time run for the operation; useful when there is grouping
SkipCols	Defaults to NULL; otherwise name the vector containing the names of columns to skip
SimpleImpute	Set to TRUE for factor level imputation of "0" and numeric imputation of -1

Value

data.table of original data plus newly created features

Author(s)

Adrian Antico

See Also

Other Feature Engineering: AutoWord2VecModeler, DummifyDT, FAST_GDL_Feature_Engineering, GDL_Feature_Engineering, ModelDataPrep, Scoring_GDL_Feature_Engineering

Examples

```
N = 25116
data <- data.table::data.table(DateTime = as.Date(Sys.time()),
                               Target = stats::filter(rnorm(N,
                                                            mean = 50,
                                                            sd = 20),
                                                      filter=rep(1,10),
                                                      circular=TRUE))
data[, temp := seq(1:N)][, DateTime := DateTime - temp][, temp := NULL]
data <- data[order(DateTime)]
data <- DT_GDL_Feature_Engineering(data,
                                   lags          = c(seq(1,5,1)),
                                   periods       = c(3,5,10,15,20,25),
                                   statsNames    = c("MA"),
                                   targets       = c("Target"),
                                   groupingVars  = NULL,
                                   sortDateName  = "DateTime",
                                   timeDiffTarget = c("Time_Gap"),
                                   timeAgg       = c("days"),
                                   WindowingLag  = 1,
                                   Type          = "Lag",
                                   Timer         = TRUE,
                                   SkipCols      = FALSE,
                                   SimpleImpute  = TRUE)
```

DummifyDT	*DummifyDT creates dummy variables for the selected columns.*

Description

DummifyDT creates dummy variables for the selected columns. Either one-hot encoding, N+1 columns for N levels, or N columns for N levels.

Usage

```
DummifyDT(data, cols, KeepBaseCols = TRUE, OneHot = TRUE)
```

Arguments

data	the data set to run the micro auc on
cols	a vector with the names of the columns you wish to dichotomize
KeepBaseCols	set to TRUE to keep the original columns used in the dichotomization process
OneHot	Set to TRUE to run one hot encoding, FALSE to generate N columns for N levels

Value

data table with new dummy variables columns and optionally removes base columns

Author(s)

Adrian Antico

See Also

Other Feature Engineering: AutoWord2VecModeler, DT_GDL_Feature_Engineering, FAST_GDL_Feature_Engineering, GDL_Feature_Engineering, ModelDataPrep, Scoring_GDL_Feature_Engineering

Examples

```
test <- data.table::data.table(Value = runif(100000),
                  FactorCol = sample(x = c(letters,
                                            LETTERS,
                                            paste0(letters,letters),
                                            paste0(LETTERS,LETTERS),
                                            paste0(letters,LETTERS),
                                            paste0(LETTERS,letters)),
                                      size = 100000,
                                      replace = TRUE))
test <- DummifyDT(data = test,
                  cols = "FactorCol",
                  KeepBaseCols = FALSE)
ncol(test)
test[, sum(FactorCol_gg)]
```

EvalPlot   *Function automatically builds calibration plots for model evalua-*
                            tion

Description

This function automatically builds calibration plots and calibration boxplots for model evaluation using regression, quantile regression, and binary and multinomial classification

Usage

```
EvalPlot(data, PredColName = c("PredictedValues"),
  ActColName = c("ActualValues"), type = c("calibration"),
  bucket = 0.05, aggrfun = function(x) base::mean(x, na.rm = TRUE))
```

Arguments

data	Data containing predicted values and actual values for comparison
PredColName	String representation of column name with predicted values from model
ActColName	String representation of column name with actual values from model
type	Calibration or boxplot - calibration aggregated data based on summary statistic; boxplot shows variation
bucket	Number of buckets to partition the space on (0,1) for evaluation
aggrfun	The statistics function used in aggregation, listed as a function

Value

Calibration plot or boxplot

Author(s)

Adrian Antico

See Also

Other Model Evaluation and Interpretation: ParDepCalPlots, RedYellowGreen, threshOptim

Examples

```
## Not run:
EvalPlot(data,
         PredColName = "predict",
         ActColName  = "target",
         type        = "calibration",
         bucket      = 0.05,
         aggrfun     = function(x) quantile(x, probs = 0.50, na.rm = TRUE))

## End(Not run)
```

FAST_GDL_Feature_Engineering

An Fast Automated Feature Engineering Function

Description

For models with target variables within the realm of the current time frame but not too far back in time, this function creates autoregressive and rolling stats from target columns and distributed lags and distributed rolling stats for independent features distributed across time. On top of that, you can also create time between instances along with their associated lags and rolling stats. This function works for data with groups and without groups.

Usage

```
FAST_GDL_Feature_Engineering(data, lags = c(1:5), periods = c(seq(10,
  50, 10)), statsFUNs = c("mean", "median", "sd", "quantile85",
  "quantile95"), statsNames = c("mean", "median", "sd", "quantile85",
  "quantile95"), targets = c("Target"),
  groupingVars = c("GroupVariable"), sortDateName = c("DateTime"),
  timeDiffTarget = NULL, timeAgg = c("hours"), WindowingLag = 1,
  Type = c("Lag"), Timer = FALSE, SkipCols = FALSE,
  SimpleImpute = TRUE, AscRowByGroup = c("temp"), RecordsKeep = 1)
```

Arguments

data	The data source you want to run the function on
lags	The list of specific lags you want to have generated
periods	The number of periods for the rolling stats

statsFUNs	List of functions for your rolling windows, such as mean, sd, min, max, quantile
statsNames	The corresponding names to append to your colnames created associated with statsFuns
targets	The column(s) in which you will build your lags and rolling stats
groupingVars	Categorical variables you will build your lags and rolling stats by
sortDateName	String name of your core date column in your transaction data
timeDiffTarget	List a name in order to create time between events with assiciated lags and rolling features
timeAgg	Unit of time to aggregate by
WindowingLag	Build moving stats off of target column(s) or one of their lags (1+)
Type	input "Lag" if you want features built on historical values; use "Lead" if you want features built on future values
Timer	Set to TRUE if you want a time run for the operation; useful when there is grouping
SkipCols	Defaults to NULL; otherwise name the vector containing the names of columns to skip
SimpleImpute	Set to TRUE for factor level imputation of "0" and numeric imputation of -1
AscRowByGroup	Required to have a column with a Row Number by group (if grouping) with 1 being the record for scoring (typically the most current in time)
RecordsKeep	List the number of records to retain (1 for last record, 2 for last 2 records, etc.)

Value

data.table of original data plus newly created features

Author(s)

Adrian Antico

See Also

Other Feature Engineering: AutoWord2VecModeler, DT_GDL_Feature_Engineering, DummifyDT, GDL_Feature_Engineering, ModelDataPrep, Scoring_GDL_Feature_Engineering

Examples

```
N = 25116
data <- data.table::data.table(DateTime = as.Date(Sys.time()),
  Target = stats::filter(rnorm(N,
                               mean = 50,
                               sd = 20),
                         filter=rep(1,10),
                         circular=TRUE))
data[, temp := seq(1:N)][, DateTime := DateTime - temp]
data <- data[order(DateTime)]
data <- FAST_GDL_Feature_Engineering(data,
                                     lags          = c(1:5),
                                     periods       = c(seq(10,50,10)),
```

```
                                statsFUNs      = c("mean",
                                                   "median",
                                                   "sd",
                                                   "quantile85",
                                                   "quantile95"),
                                statsNames     = c("mean",
                                                   "median",
                                                   "sd",
                                                   "quantile85",
                                                   "quantile95"),
                                targets        = c("Target"),
                                groupingVars   = NULL,
                                sortDateName   = "DateTime",
                                timeDiffTarget = c("Time_Gap"),
                                timeAgg        = "days",
                                WindowingLag   = 1,
                                Type           = "Lag",
                                Timer          = TRUE,
                                SkipCols       = FALSE,
                                SimpleImpute   = TRUE,
                                AscRowByGroup  = "temp")
```

GDL_Feature_Engineering

An Automated Feature Engineering Function

Description

Builds autoregressive and rolling stats from target columns and distributed lags and distributed rolling stats for independent features distributed across time. On top of that, you can also create time between instances along with their associated lags and rolling stats. This function works for data with groups and without groups.

Usage

```
GDL_Feature_Engineering(data, lags = c(seq(1, 5, 1)), periods = c(3, 5,
  10, 15, 20, 25), statsFUNs = c(function(x) quantile(x, probs = 0.1,
  na.rm = TRUE), function(x) quantile(x, probs = 0.9, na.rm = TRUE),
  function(x) base::mean(x, na.rm = TRUE), function(x) sd(x, na.rm = TRUE),
  function(x) quantile(x, probs = 0.25, na.rm = TRUE), function(x)
  quantile(x, probs = 0.75, na.rm = TRUE)), statsNames = c("q10", "q90",
  "mean", "sd", "q25", "q75"), targets = c("qty"),
  groupingVars = c("Group1", "Group2"), sortDateName = c("date"),
  timeDiffTarget = c("TimeDiffName"), timeAgg = c("days"),
  WindowingLag = 0, Type = c("Lag"), Timer = TRUE, SkipCols = NULL,
  SimpleImpute = TRUE)
```

Arguments

data	The data source you want to run the function on
lags	The list of specific lags you want to have generated
periods	The number of periods to use for rolling stats

statsFUNs	List of functions for your rolling windows, such as mean, sd, min, max, quantile
statsNames	The corresponding names to append to your colnames created associated with statsFuns
targets	The column(s) in which you will build your lags and rolling stats
groupingVars	Categorical variables you will build your lags and rolling stats by
sortDateName	String name of your core date column in your transaction data
timeDiffTarget	List a name in order to create time between events with assiciated lags and rolling features
timeAgg	Unit of time to aggregate by
WindowingLag	Build moving stats off of target column(s) or one of their lags (1+)
Type	input "Lag" if you want features built on historical values; use "Lead" if you want features built on future values
Timer	Set to TRUE if you want a time run for the operation; useful when there is grouping
SkipCols	Defaults to NULL; otherwise name the vector containing the names of columns to skip
SimpleImpute	Set to TRUE for factor level imputation of "0" and numeric imputation of -1

Value

data.table of original data plus newly created features

Author(s)

Adrian Antico

See Also

Other Feature Engineering: AutoWord2VecModeler, DT_GDL_Feature_Engineering, DummifyDT, FAST_GDL_Feature_Engineering, ModelDataPrep, Scoring_GDL_Feature_Engineering

Examples

```
N = 25116
data <- data.table::data.table(DateTime = as.Date(Sys.time()),
  Target = stats::filter(rnorm(N,
                               mean = 50,
                               sd = 20),
                         filter=rep(1,10),
                         circular=TRUE))
data[, temp := seq(1:N)][, DateTime := DateTime - temp][, temp := NULL]
data <- data[order(DateTime)]
data <- GDL_Feature_Engineering(data,
         lags             = c(seq(1,5,1)),
         periods          = c(3,5,10,15,20,25),
         statsFUNs        = c(function(x) quantile(x, probs = 0.20, na.rm = TRUE),
                              function(x) quantile(x, probs = 0.80, na.rm = TRUE),
                              function(x) mean(x, na.rm = TRUE),
                              function(x) sd(x, na.rm = TRUE),
                              function(x) quantile(x, probs = 0.10, na.rm = TRUE),
```

```
                             function(x) quantile(x, probs = 0.90, na.rm = TRUE)),
             statsNames    = c("min","max","mean","sd","q20","q80"),
             targets       = c("Target"),
             groupingVars  = NULL,
             sortDateName  = "DateTime",
             timeDiffTarget = c("Time_Gap"),
             timeAgg       = "days",
             WindowingLag  = 1,
             Type          = "Lag",
             Timer         = TRUE,
             SkipCols      = FALSE,
             SimpleImpute  = TRUE)
```

GenTSAnomVars	*GenTSAnomVars is an automated z-score anomaly detection via GLM-like procedure*

Description

GenTSAnomVars is an automated z-score anomaly detection via GLM-like procedure. Data is z-scaled and grouped by factors and time periods to determine which points are above and below the control limits in a cumulative time fashion. Then a cumulative rate is created as the final variable. Set KeepAllCols to FALSE to utilize the intermediate features to create rolling stats from them.

Usage

```
GenTSAnomVars(data, ValueCol = "Value", GroupVar1 = "SKU",
  GroupVar2 = NULL, DateVar = "DATE", High = 1.96, Low = -1.96,
  KeepAllCols = FALSE, DataScaled = TRUE)
```

Arguments

data	the source residuals data.table
ValueCol	the numeric column to run anomaly detection over
GroupVar1	this is a group by variable
GroupVar2	this is another group by variable
DateVar	this is a time variable for grouping
High	this is the threshold on the high end
Low	this is the threshold on the low end
KeepAllCols	set to TRUE to remove the intermediate features
DataScaled	set to TRUE if you already scaled your data

Value

The original data.table with the added columns merged in

Author(s)

Adrian Antico

See Also

Other Unsupervised Learning: AutoKMeans, ResidualOutliers

Examples

```
data <- data.table::data.table(DateTime = as.Date(Sys.time()),
  Target = stats::filter(rnorm(1000,
                               mean = 50,
                               sd = 20),
                      filter=rep(1,10),
                      circular=TRUE))
data[, temp := seq(1:10000)][, DateTime := DateTime - temp][, temp := NULL]
data <- data[order(DateTime)]
x <- data.table::as.data.table(sde::GBM(N=10000)*1000)
data[, predicted := x[-1,]]
stuff    <- GenTSAnomVars(data,
                          ValueCol    = "Target",
                          GroupVar1   = NULL,
                          GroupVar2   = NULL,
                          DateVar     = "DateTime",
                          High        = 1.96,
                          Low         = -1.96,
                          KeepAllCols = TRUE,
                          DataScaled  = FALSE)
```

ModelDataPrep   *Final Data Preparation Function*

Description

This function replaces inf values with NA, converts characters to factors, and imputes with constants

Usage

```
ModelDataPrep(data, Impute = TRUE, CharToFactor = TRUE,
  MissFactor = "0", MissNum = -1)
```

Arguments

data	This is your source data you'd like to modify
Impute	Defaults to TRUE which tells the function to impute the data
CharToFactor	Defaults to TRUE which tells the function to convert characters to factors
MissFactor	Supply the value to impute missing factor levels
MissNum	Supply the value to impute missing numeric values

Value

Returns the original data table with corrected values

Author(s)

Adrian Antico

See Also

Other Feature Engineering: AutoWord2VecModeler, DT_GDL_Feature_Engineering, DummifyDT, FAST_GDL_Feature_Engineering, GDL_Feature_Engineering, Scoring_GDL_Feature_Engineering

Examples

```
data <- data.table::data.table(Value = runif(100000),
                               FactorCol = as.character(sample(x = c(letters,
                                                                     LETTERS,
                                                                paste0(letters,letters),
                                                                paste0(LETTERS,LETTERS),
                                                                paste0(letters,LETTERS),
                                                               paste0(LETTERS,letters)),
                                                               size = 100000,
                                                               replace = TRUE)))
data <- ModelDataPrep(data,
                      Impute = TRUE,
                      CharToFactor = TRUE,
                      MissFactor = "0",
                      MissNum    = -1)
```

---

multiplot	*Multiplot is a function for combining multiple plots*

Description

Sick of copying this one into your code? Well, not anymore.

Usage

```
multiplot(..., plotlist = NULL, cols = 2, layout = NULL)
```

Arguments

...	Passthrough arguments
plotlist	This is the list of your charts
cols	This is the number of columns in your multiplot
layout	Leave NULL

Value

Multiple ggplots on a single image

Author(s)

Adrian Antico

See Also

Other Misc: AutoH20TextPrepScoring, AutoWordFreq, ChartTheme, PrintObjectsSize, RecomDataCreate, RemixTheme, SimpleCap, percRank, tempDatesFun, tokenizeH20

Examples

```
## Not run:
multiplot(plotlist = list(p1,p2,p3,p4), cols = 2)

## End(Not run)
```

ParDepCalPlots	*Function automatically builds partial dependence calibration plots for model evaluation*

Description

This function automatically builds partial dependence calibration plots and partial dependence calibration boxplots for model evaluation using regression, quantile regression, and binary and multinomial classification

Usage

```
ParDepCalPlots(data, PredColName = c("PredictedValues"),
  ActColName = c("ActualValues"),
  IndepVar = c("Independent_Variable_Name"), type = c("calibration"),
  bucket = 0.05, FactLevels = 10, Function = function(x)
  base::mean(x, na.rm = TRUE))
```

Arguments

data	Data containing predicted values and actual values for comparison
PredColName	String representation of the column name with predicted values from model
ActColName	String representation of the column name with actual values from model
IndepVar	String representation of the column name with the independent variable of choice
type	Calibration or boxplot - calibration aggregated data based on summary statistic; boxplot shows variation
bucket	Number of buckets to partition the space on (0,1) for evaluation
FactLevels	The number of levels to show on the chart (1. Levels are chosen based on frequency; 2. all other levels grouped and labeled as "Other")
Function	Supply the function you wish to use for aggregation.

Value

Partial dependence calibration plot or boxplot

Author(s)

Adrian Antico

See Also

Other Model Evaluation and Interpretation: EvalPlot, RedYellowGreen, threshOptim

Examples

```
## Not run:
ParDepCalPlots(data,
               PredColName = "predict",
               ActColName  = "target",
               IndepVar    = "Independent_Variable",
               type        = "boxplot",
               bucket      = 0.05,
               FactLevels  = 10,
               Function    = function(x) mean(x, na.rm = TRUE))

## End(Not run)
```

---

percRank   *Percentile rank function*

Description

This function computes percentile ranks for each row in your data like Excel's PER-CENT_RANK

Usage

```
percRank(x)
```

Arguments

x   X is your variable of interest

Value

vector of percentile ranks

Author(s)

Adrian Antico

See Also

Other Misc: AutoH20TextPrepScoring, AutoWordFreq, ChartTheme, PrintObjectsSize, RecomDataCreate, RemixTheme, SimpleCap, multiplot, tempDatesFun, tokenizeH20

Examples

```
## Not run:
percRank(x)

## End(Not run)
```

PrintObjectsSize   *PrintObjectsSize prints out the top N objects and their associated sizes, sorted by size*

Description

PrintObjectsSize prints out the top N objects and their associated sizes, sorted by size

Usage

```
PrintObjectsSize(N = 10)
```

Arguments

N   The number of objects to display

Value

The objects in your environment and their sizes

Author(s)

Adrian Antico

See Also

Other Misc: AutoH20TextPrepScoring, AutoWordFreq, ChartTheme, RecomDataCreate, RemixTheme, SimpleCap, multiplot, percRank, tempDatesFun, tokenizeH20

Examples

```
## Not run:
PrintObjectsSize(N = 10)

## End(Not run)
```

RecomDataCreate   *Convert transactional data.table to a binary ratings matrix*

Description

Convert transactional data.table to a binary ratings matrix

Usage

```
RecomDataCreate(data, EntityColName = "CustomerID",
  ProductColName = "StockCode", MetricColName = "TotalSales")
```

Arguments

`data`	This is your transactional data.table. Must include an Entity (typically customer), ProductCode (such as SKU), and a sales metric (such as total sales).
`EntityColName`	This is the column name in quotes that represents the column name for the Entity, such as customer
`ProductColName`	This is the column name in quotes that represents the column name for the product, such as SKU
`MetricColName`	This is the column name in quotes that represents the column name for the metric, such as total sales

Value

A BinaryRatingsMatrix

Author(s)

Adrian Antico and Douglas Pestana

See Also

Other Misc: AutoH20TextPrepScoring, AutoWordFreq, ChartTheme, PrintObjectsSize, RemixTheme, SimpleCap, multiplot, percRank, tempDatesFun, tokenizeH20

Examples

```
## Not run:
BinaryRatingsMatrix <- RecomDataCreate(data,
                                    EntityColName = "CustomerID",
                                    ProductColName = "StockCode",
                                    MetricColName = "TotalSales")


## End(Not run)
```

RedYellowGreen   *RedYellowGreen is for determining the optimal thresholds for bi-
                            nary classification when do-nothing is an option*

Description

This function will find the optimal thresholds for applying the main label and for finding the optimal range for doing nothing when you can quantity the cost of doing nothing

Usage

```
RedYellowGreen(data, PredictColNumber = 2, ActualColNumber = 1,
  TruePositiveCost = 0, TrueNegativeCost = 0,
  FalsePositiveCost = -10, FalseNegativeCost = -50, MidTierCost = -2,
  Cores = 8, Precision = 0.01)
```

Arguments

`data`	data is the data table with your predicted and actual values from a classification model
`PredictColNumber`	The column number where the actual target variable is located (in binary form)
`ActualColNumber`	The column number where the predicted values are located
`TruePositiveCost`	This is the utility for generating a true positive prediction
`TrueNegativeCost`	This is the utility for generating a true negative prediction
`FalsePositiveCost`	This is the cost of generating a false positive prediction
`FalseNegativeCost`	This is the cost of generating a false negative prediction
`MidTierCost`	This is the cost of doing nothing (or whatever it means to not classify in your case)
`Cores`	Number of cores on your machine
`Precision`	Set the decimal number to increment by between 0 and 1

Value

A data table with all evaluated strategies, parameters, and utilities, along with a 3d scatterplot of the results

Author(s)

Adrian Antico

See Also

Other Model Evaluation and Interpretation: `EvalPlot`, `ParDepCalPlots`, `threshOptim`

Examples

```
## Not run:
data <- RedYellowGreen(data,
                      PredictColNumber  = 1,
                      ActualColNumber   = 2,
                      TruePositiveCost  = 0,
                      TrueNegativeCost  = 0,
                      FalsePositiveCost = -1,
                      FalseNegativeCost = -2,
                      MidTierCost       = -0.5)

## End(Not run)
```

RemixTheme   *RemixTheme function is a ggplot theme generator for ggplots*

Description

This function adds the Remix Theme to ggplots

Usage

```
RemixTheme()
```

Value

An object to pass along to ggplot objects following the "+" sign

Author(s)

Douglas Pestana

See Also

Other Misc: AutoH20TextPrepScoring, AutoWordFreq, ChartTheme, PrintObjectsSize, RecomDataCreate, SimpleCap, multiplot, percRank, tempDatesFun, tokenizeH20

Examples

```
data <- data.table::data.table(DateTime = as.Date(Sys.time()),
  Target = stats::filter(rnorm(1000,
                                mean = 50,
                                sd = 20),
                        filter=rep(1,10),
                        circular=TRUE))
data[, temp := seq(1:1000)][, DateTime := DateTime - temp][, temp := NULL]
data <- data[order(DateTime)]
p <- ggplot2::ggplot(data, ggplot2::aes(x = DateTime, y = Target)) + ggplot2::geom_line()
p <- p + RemixTheme()
```

ResidualOutliers   *ResidualOutliers is an automated time series outlier detection function*

Description

ResidualOutliers is an automated time series outlier detection function that utilizes tsoutliers and auto.arima.

Usage

```
ResidualOutliers(data, maxN = 5, cvar = 4)
```

Arguments

data	the source residuals data.table
maxN	the largest lag or moving average (seasonal too) values for the arima fit
cvar	the t-stat value for tsoutliers

Value

A data.table with outliers, the arima model, and residuals from the arima fit

Author(s)

Adrian Antico

See Also

Other Unsupervised Learning: AutoKMeans, GenTSAnomVars

Examples

```
data <- data.table::data.table(a = seq(0,10000,1),
                   predicted = sde::GBM(N=10000)*1000)[-1,]
data <- data.table::data.table(a = seq(1,10000,1),
                   predicted = sde::rcCIR(n=10000,
                                          Dt=0.1,
                                          x0=1,
                                          theta=c(6,2,2)))
data <- data.table::data.table(a = seq(1,10000,1),
                   predicted = sde::rsOU(n=10000,
                                          theta=c(0,2,1)))
stuff    <- ResidualOutliers(data = data, maxN = 5, cvar = 4)
data     <- stuff[[1]]
model    <- stuff[[2]]
resid    <- stuff[[3]]
outliers <- data[type != "<NA>"]
```

Scoring_GDL_Feature_Engineering

An Automated Scoring Feature Engineering Function

Description

For scoring purposes (brings back a single row by group), this function creates autoregressive and rolling stats from target columns and distributed lags and distributed rolling stats for independent features distributed across time. On top of that, you can also create time between instances along with their associated lags and rolling stats. This function works for data with groups and without groups.

Usage

```
Scoring_GDL_Feature_Engineering(data, lags = c(1:6, 12, seq(24, 168,
  24)), periods = c(6, 12, 24, 72, 168, 720, 4320, 8640),
  statsFUNs = c(function(x) base::mean(x, na.rm = TRUE), function(x)
  base::sd(x, na.rm = TRUE)), statsNames = c("mean", "sd"),
  targets = c("Target"), groupingVars = c("GroupVariable"),
  sortDateName = c("DateTime"), timeDiffTarget = c("Time_Gap"),
  timeAgg = c("days"), WindowingLag = 1, Type = c("Lag"),
  Timer = FALSE, SkipCols = FALSE, SimpleImpute = TRUE,
  AscRowByGroup = c("temp"), RecordsKeep = 1)
```

Arguments

`data`	The data source you want to run the function on
`lags`	The list of specific lags you want to have generated
`periods`	The number of periods in the rolling statistics
`statsFUNs`	List of functions for your rolling windows, such as mean, sd, min, max, quantile
`statsNames`	The corresponding names to append to your colnames created associated with statsFuns
`targets`	The column(s) in which you will build your lags and rolling stats
`groupingVars`	Categorical variables you will build your lags and rolling stats by
`sortDateName`	String name of your core date column in your transaction data
`timeDiffTarget`	List a name in order to create time between events with assiciated lags and rolling features
`timeAgg`	Unit of time to aggregate by
`WindowingLag`	Build moving stats off of target column(s) or one of their lags (1+)
`Type`	input "Lag" if you want features built on historical values; use "Lead" if you want features built on future values
`Timer`	Set to TRUE if you want a time run for the operation; useful when there is grouping
`SkipCols`	Defaults to NULL; otherwise name the vector containing the names of columns to skip
`SimpleImpute`	Set to TRUE for factor level imputation of "0" and numeric imputation of -1
`AscRowByGroup`	Required to have a column with a Row Number by group (if grouping) with 1 being the record for scoring (typically the most current in time)
`RecordsKeep`	List the number of records to retain (1 for last record, 2 for last 2 records, etc.)

Value

data.table of original data plus newly created features

Author(s)

Adrian Antico

See Also

Other Feature Engineering: AutoWord2VecModeler, DT_GDL_Feature_Engineering, DummifyDT, FAST_GDL_Feature_Engineering, GDL_Feature_Engineering, ModelDataPrep

Examples

```
N = 25116
data1 <- data.table::data.table(DateTime = as.Date(Sys.time()),
                                Target = stats::filter(rnorm(N,
                                                             mean = 50,
                                                             sd = 20),
                                                       filter=rep(1,10),
                                                       circular=TRUE))
data1[, temp := seq(1:N)][, DateTime := DateTime - temp]
data1 <- data1[order(DateTime)]
data1 <- Scoring_GDL_Feature_Engineering(data1,
                                lags          = c(seq(1,5,1)),
                                periods       = c(3,5,10,15,20,25),
                              statsFUNs     = c(function(x) mean(x,na.rm = TRUE)),
                                statsNames    = c("MA"),
                                targets       = c("Target"),
                                groupingVars  = NULL,
                                sortDateName  = c("DateTime"),
                                timeDiffTarget = c("Time_Gap"),
                                timeAgg       = "days",
                                WindowingLag  = 1,
                                Type          = "Lag",
                                Timer         = TRUE,
                                SkipCols      = FALSE,
                                SimpleImpute  = TRUE,
                                AscRowByGroup = "temp",
                                RecordsKeep   = 1)
```

SimpleCap	SimpleCap function is for capitalizing the first letter of words

---

Description

SimpleCap function is for capitalizing the first letter of words (need I say more?)

Usage

```
SimpleCap(x)
```

Arguments

x   Column of interest

Value

An object to pass along to ggplot objects following the "+" sign

Author(s)

Adrian Antico

See Also

Other Misc: AutoH20TextPrepScoring, AutoWordFreq, ChartTheme, PrintObjectsSize, RecomDataCreate, RemixTheme, multiplot, percRank, tempDatesFun, tokenizeH20

Examples

```
x <- "adrian"
x <- SimpleCap(x)
```

---

tempDatesFun	*tempDatesFun Convert Excel datetime char columns to Date columns*

Description

tempDatesFun takes the Excel datetime column, which imports as character, and converts it into a date type

Usage

```
tempDatesFun(x)
```

Arguments

x   The column of interest

Value

An object to pass along to ggplot objects following the "+" sign

Author(s)

Adrian Antico

See Also

Other Misc: AutoH20TextPrepScoring, AutoWordFreq, ChartTheme, PrintObjectsSize, RecomDataCreate, RemixTheme, SimpleCap, multiplot, percRank, tokenizeH20

Examples

```
## Not run:
Cdata[, DAY_DATE := tempDatesFun(DAY_DATE)]
Cdata[, DAY_DATE := base::as.Date(DAY_DATE, "%m/%d/%Y")]

## End(Not run)
```

---

threshOptim   *Utility maximizing thresholds for binary classification*

Description

This function will return the utility maximizing threshold for future predictions along with the data generated to estimate the threshold

Usage

```
threshOptim(data, actTar = "target", predTar = "p1", tpProfit = 0,
  tnProfit = 0, fpProfit = -1, fnProfit = -2)
```

Arguments

data	data is the data table you are building the modeling on
actTar	The column name where the actual target variable is located (in binary form)
predTar	The column name where the predicted values are located
tpProfit	This is the utility for generating a true positive prediction
tnProfit	This is the utility for generating a true negative prediction
fpProfit	This is the cost of generating a false positive prediction
fnProfit	This is the cost of generating a false negative prediction

Value

Optimal threshold and corresponding utilities for the range of thresholds tested

Author(s)

Adrian Antico

See Also

Other Model Evaluation and Interpretation: EvalPlot, ParDepCalPlots, RedYellowGreen

Examples

```
## Not run:
data <- threshOptim(data    = data,
                    actTar   = "target",
                    predTar  = "p1",
                    tpProfit = 0,
                    tnProfit = 0,
                    fpProfit = -1,
                    fnProfit = -2)
optimalThreshold <- data[[1]]
allResults       <- data[[2]]

## End(Not run)
```

tokenizeH20	*For NLP work*

Description

This function tokenizes data

Usage

```
tokenizeH20(data)
```

Arguments

data   The text data

Author(s)

Adrian Antico

See Also

Other Misc: AutoH20TextPrepScoring, AutoWordFreq, ChartTheme, PrintObjectsSize, RecomDataCreate, RemixTheme, SimpleCap, multiplot, percRank, tempDatesFun

Index